

Title : Electric Vehicle Data Analysis Assignment

Name : Mamta Bhagde

Date : 10-10-2025

Course: Data Analysis

Introduction : The dataset contains information on Battery Electric Vehicles (BEVs) and Plug-in Hybrid Electric Vehicles (PHEVs). The main aim of this study is to analyze electric vehicle (EV), understand EV adoption patterns across different regions and vehicle types, and build predictive models to estimate the electric range of vehicles. By examining vehicle characteristics, pricing, and incentives, this analysis provides a comprehensive view of EV trends.

Section 1:Data Cleaning Questions

1.How many missing values exist in the dataset, and in which columns?

```
import pandas as pd
df = pd.read_csv(r"D:\DA\python\Electric_Vehicle_Population_Data
(3).csv")

# Count missing values
missing_values = df.isnull().sum()
missing_report = pd.DataFrame({
    'Missing Values': missing_values})
# Print the missing values
print(missing_report)
```

	Missing Values
VIN (1-10)	0
County	10
City	10
State	0
Postal Code	10
Model Year	0
Make	0
Model	0
Electric Vehicle Type	0
Clean Alternative Fuel Vehicle (CAFV) Eligibility	0
Electric Range	3
Base MSRP	3

Legislative District	628
DOL Vehicle ID	0
Vehicle Location	18
Electric Utility	10
2020 Census Tract	10

- How should missing or zero values in the Base MSRP and Electric Range columns be handled?

```
# Missing values = the cells have no data.
# Zero values = For Base MSRP or Electric Range, a value of zero
# doesn't make sense – cars don't cost $0 and EVs can't have 0 miles of
# range

# Remove rows where Base MSRP or Electric Range is missing or zero
df_clean = df[(df['Base MSRP'] > 0) & (df['Electric Range'] > 0)]

# Replace zeros or NaN in Base MSRP
median_msrp = df['Base MSRP'][df['Base MSRP'] > 0].median()
df['Base MSRP'] = df['Base MSRP'].replace(0, median_msrp)
df['Base MSRP'] = df['Base MSRP'].fillna(median_msrp)

# Replace zeros or NaN in Electric Range
median_range = df['Electric Range'][df['Electric Range'] > 0].median()
df['Electric Range'] = df['Electric Range'].replace(0, median_range)
df['Electric Range'] = df['Electric Range'].fillna(median_range)

print(df['Base MSRP'].describe())
print(df['Electric Range'].describe())

count    261698.000000
mean      59866.481612
std       3005.824570
min       31950.000000
25%      59900.000000
50%      59900.000000
75%      59900.000000
max      845000.000000
Name: Base MSRP, dtype: float64
count    261698.000000
mean      75.198794
std       66.975490
min       1.000000
25%      53.000000
50%      53.000000
75%      53.000000
max      337.000000
Name: Electric Range, dtype: float64
```

- Are there duplicate records in the dataset? If so, how should they be managed?

```

# Checking for duplicate rows
duplicates = df.duplicated()

# total duplicates
total_duplicates = duplicates.sum()
print("Total duplicate rows:", total_duplicates)

Total duplicate rows: 0

# Create a new column 'VIN_anon' with unique numeric IDs
df['VIN_anon'] = range(1, len(df)+1)

# Drop the original VIN if needed
df = df.drop(columns=['VIN (1-10)'])
df.head()

      County          City State Postal Code Model Year Make
Model \
0    Yakima        Yakima   WA    98902.0  2013  TOYOTA
PRIUS
1    Kitsap  Port Orchard   WA    98366.0  2025    FORD
ESCAPE
2    Kitsap       Kingston   WA    98346.0  2024  MAZDA
CX-90
3  Thurston       Olympia   WA    98501.0  2023  TESLA
MODEL Y
4  Thurston       Rainier   WA    98576.0  2019  TESLA
MODEL 3

                    Electric Vehicle Type \
0  Plug-in Hybrid Electric Vehicle (PHEV)
1  Plug-in Hybrid Electric Vehicle (PHEV)
2  Plug-in Hybrid Electric Vehicle (PHEV)
3           Battery Electric Vehicle (BEV)
4           Battery Electric Vehicle (BEV)

      Clean Alternative Fuel Vehicle (CAFV) Eligibility Electric
Range \
0            Not eligible due to low battery range             6.0
1            Clean Alternative Fuel Vehicle Eligible          37.0
2            Not eligible due to low battery range          26.0
3  Eligibility unknown as battery range has not b...          53.0
4            Clean Alternative Fuel Vehicle Eligible         220.0

      Base MSRP  Legislative District  DOL Vehicle ID \
0     59900.0           15.0        165252538

```

1	59900.0	26.0	278572521
2	59900.0	23.0	275123642
3	59900.0	35.0	249569323
4	59900.0	20.0	283135107

Vehicle Location Electric Utility 2020 Census
Tract \

0	POINT (-120.51904 46.59783)	PACIFICORP
5.307700e+10		
1	POINT (-122.63847 47.54103)	PUGET SOUND ENERGY INC
5.303509e+10		
2	POINT (-122.4977 47.79802)	PUGET SOUND ENERGY INC
5.303509e+10		
3	POINT (-122.89165 47.03954)	PUGET SOUND ENERGY INC
5.306701e+10		
4	POINT (-122.68993 46.88897)	PUGET SOUND ENERGY INC
5.306701e+10		

VIN_anon

0	1
1	2
2	3
3	4
4	5

Vehicle Location is stored as GPS coordinates (longitude, latitude).
They may have issues like: Missing values, Wrong formatting (e.g.,
strings instead of numbers), Hard to read in analysis or maps

- How can VINs be anonymized while maintaining uniqueness?

```
import hashlib

df = pd.read_csv(r"D:\DA\python\Electric_Vehicle_Population_Data
(3).csv")

-----
NameError                               Traceback (most recent call
last)
Cell In[3], line 1
----> 1 df = pd.read_csv(r"D:\DA\python\
Electric_Vehicle_Population_Data (3).csv")

NameError: name 'pd' is not defined

import pandas as pd
import hashlib

df = pd.read_csv(r"D:\DA\python\Electric_Vehicle_Population_Data
(3).csv")
```

```
df.head()
```

Make \	VIN (1-10)	County	City	State	Postal Code	Model Year
TOYOTA	JTDKN3DP2D	Yakima	Yakima	WA	98902.0	2013
FORD	1FMCU0E1XS	Kitsap	Port Orchard	WA	98366.0	2025
MAZDA	JM3KKBHA9R	Kitsap	Kingston	WA	98346.0	2024
TESLA	7SAYGDEE8P	Thurston	Olympia	WA	98501.0	2023
TESLA	5YJ3E1EB5K	Thurston	Rainier	WA	98576.0	2019

Model \	Electric Vehicle Type \
PRIUS	Plug-in Hybrid Electric Vehicle (PHEV)
ESCAPE	Plug-in Hybrid Electric Vehicle (PHEV)
CX-90	Plug-in Hybrid Electric Vehicle (PHEV)
MODEL Y	Battery Electric Vehicle (BEV)
MODEL 3	Battery Electric Vehicle (BEV)

Range \	Clean Alternative Fuel Vehicle (CAFV) Eligibility	Electric
0	Not eligible due to low battery range	6.0
1	Clean Alternative Fuel Vehicle Eligible	37.0
2	Not eligible due to low battery range	26.0
3	Eligibility unknown as battery range has not b...	0.0
4	Clean Alternative Fuel Vehicle Eligible	220.0

Base MSRP \	Legislative District	DOL	Vehicle ID \
0.0	15.0	165252538	
0.0	26.0	278572521	
0.0	23.0	275123642	
0.0	35.0	249569323	
0.0	20.0	283135107	

Tract	Vehicle Location	Electric Utility	2020 Census
0	POINT (-120.51904 46.59783)	PACIFICORP	
5.307700e+10			
1	POINT (-122.63847 47.54103)	PUGET SOUND ENERGY INC	
5.303509e+10			
2	POINT (-122.4977 47.79802)	PUGET SOUND ENERGY INC	
5.303509e+10			

```

3 POINT (-122.89165 47.03954) PUGET SOUND ENERGY INC
5.306701e+10
4 POINT (-122.68993 46.88897) PUGET SOUND ENERGY INC
5.306701e+10

# Function to anonymize VIN
def anonymize_vin(vin):
    return hashlib.sha256(vin.encode()).hexdigest()

# Apply to your VIN column (replace 'VIN' with your actual column name if different)
df['Anon_VIN'] = df['VIN (1-10)'].apply(anonymize_vin)
df.head()

      VIN (1-10)      County          City State Postal Code Model Year
Make \
0   JTDKN3DP2D      Yakima      Yakima  WA     98902.0  2013
TOYOTA
1   1FMCU0E1XS      Kitsap  Port Orchard  WA     98366.0  2025
FORD
2   JM3KKBHA9R      Kitsap      Kingston  WA     98346.0  2024
MAZDA
3   7SAYGDEE8P  Thurston      Olympia  WA     98501.0  2023
TESLA
4   5YJ3E1EB5K  Thurston      Rainier  WA     98576.0  2019
TESLA

      Model          Electric Vehicle Type \
0   PRIUS  Plug-in Hybrid Electric Vehicle (PHEV)
1   ESCAPE  Plug-in Hybrid Electric Vehicle (PHEV)
2   CX-90  Plug-in Hybrid Electric Vehicle (PHEV)
3   MODEL Y      Battery Electric Vehicle (BEV)
4   MODEL 3      Battery Electric Vehicle (BEV)

      Clean Alternative Fuel Vehicle (CAFV) Eligibility  Electric
Range \
0           Not eligible due to low battery range       6.0
1           Clean Alternative Fuel Vehicle Eligible    37.0
2           Not eligible due to low battery range     26.0
3   Eligibility unknown as battery range has not b...    0.0
4           Clean Alternative Fuel Vehicle Eligible   220.0

      Base MSRP  Legislative District  DOL Vehicle ID \
0        0.0          15.0        165252538
1        0.0          26.0        278572521
2        0.0          23.0        275123642

```

```

3      0.0          35.0    249569323
4      0.0          20.0    283135107

              Vehicle Location      Electric Utility 2020 Census
Tract \
0  POINT (-120.51904 46.59783)          PACIFICORP
5.307700e+10
1  POINT (-122.63847 47.54103)  PUGET SOUND ENERGY INC
5.303509e+10
2  POINT (-122.4977 47.79802)  PUGET SOUND ENERGY INC
5.303509e+10
3  POINT (-122.89165 47.03954)  PUGET SOUND ENERGY INC
5.306701e+10
4  POINT (-122.68993 46.88897)  PUGET SOUND ENERGY INC
5.306701e+10

                  Anon_VIN
0  bf01895762f04150a8ff5b0210e4d1c199986b50f45bcb...
1  a720d326091898dfa57b91cfa7466fe461b99a14f6bc78...
2  ef506f78a5a27e7e7582fd6924e13b4bfba05984680a2...
3  fb3f4d8c8632615cdf99cc78f0f8e21e1e97d1e30d6dd0...
4  5fb1eb0d5a655b4eada221a1fa28fa1c5d7fbc960c0a97...

# See how the column looks
print(df['Vehicle Location'].head(20))

0  POINT (-120.51904 46.59783)
1  POINT (-122.63847 47.54103)
2  POINT (-122.4977 47.79802)
3  POINT (-122.89165 47.03954)
4  POINT (-122.68993 46.88897)
5  POINT (-122.1389 47.87115)
6  POINT (-122.70348 47.52028)
7  POINT (-122.37265 48.24159)
8  POINT (-122.30866 47.57874)
9  POINT (-122.92333 47.03779)
10  POINT (-122.69275 47.65171)
11  POINT (-122.2066 47.67887)
12  POINT (-122.90787 46.9461)
13  POINT (-122.06402 48.01497)
14  POINT (-122.35029 47.71871)
15  POINT (-122.16335 47.53505)
16  POINT (-122.24369 47.75892)
17  POINT (-122.20563 47.76144)
18  POINT (-122.68993 46.88897)
19  POINT (-122.18637 47.89251)

Name: Vehicle Location, dtype: object

# 5. How can Vehicle Location (GPS coordinates) be cleaned or converted for better readability?

```

```

import pandas as pd

# Load CSV
df = pd.read_csv(r"D:\DA\python\Electric_Vehicle_Population_Data
(3).csv")

# Get first 5 rows
df.head()

    VIN (1-10)      County          City State Postal Code Model Year
Make \
0  JTDKN3DP2D      Yakima        Yakima WA     98902.0  2013
TOYOTA
1  1FMCU0E1XS      Kitsap Port Orchard WA     98366.0  2025
FORD
2  JM3KKBHA9R      Kitsap Kingston WA     98346.0  2024
MAZDA
3  7SAYGDEE8P      Thurston Olympia WA     98501.0  2023
TESLA
4  5YJ3E1EB5K      Thurston Rainier WA     98576.0  2019
TESLA

      Model           Electric Vehicle Type \
0    PRIUS  Plug-in Hybrid Electric Vehicle (PHEV)
1  ESCAPE  Plug-in Hybrid Electric Vehicle (PHEV)
2   CX-90  Plug-in Hybrid Electric Vehicle (PHEV)
3  MODEL Y       Battery Electric Vehicle (BEV)
4  MODEL 3       Battery Electric Vehicle (BEV)

      Clean Alternative Fuel Vehicle (CAFV) Eligibility Electric
Range \
0            Not eligible due to low battery range  6.0
1            Clean Alternative Fuel Vehicle Eligible 37.0
2            Not eligible due to low battery range  26.0
3  Eligibility unknown as battery range has not b...  0.0
4            Clean Alternative Fuel Vehicle Eligible 220.0

    Base MSRP Legislative District DOL Vehicle ID \
0      0.0          15.0  165252538
1      0.0          26.0  278572521
2      0.0          23.0  275123642
3      0.0          35.0  249569323
4      0.0          20.0  283135107

      Vehicle Location      Electric Utility 2020 Census
Tract

```

```

0 POINT (-120.51904 46.59783) PACIFICORP
5.307700e+10
1 POINT (-122.63847 47.54103) PUGET SOUND ENERGY INC
5.303509e+10
2 POINT (-122.4977 47.79802) PUGET SOUND ENERGY INC
5.303509e+10
3 POINT (-122.89165 47.03954) PUGET SOUND ENERGY INC
5.306701e+10
4 POINT (-122.68993 46.88897) PUGET SOUND ENERGY INC
5.306701e+10

# Remove invalid entries
df['Vehicle Location'] = df['Vehicle
Location'].astype(str).str.strip()
df = df[df['Vehicle Location'].notna()] # Remove missing
values
df = df[df['Vehicle Location'] != ''] # Remove empty
strings
df = df[df['Vehicle Location'].str.lower() != 'nan'] # Remove string
'nan'

# Only keep rows containing a comma (expected format: "lat, lon")
df_valid = df[df['Vehicle Location'].str.contains(',') ,
na=False)].copy()

# If no valid rows exist, we stop here
if df_valid.empty:
    print("No valid GPS coordinates found.")

No valid GPS coordinates found.

if not df_valid.empty:
    # Split by comma
    split_coords = df_valid['Vehicle Location'].str.split(',', n=1,
expand=True)

    # Make sure split worked
    if split_coords.shape[1] >= 2:
        split_coords.columns = ['Latitude', 'Longitude']

        # Convert to float safely
        def to_float(x):
            try:
                return float(str(x).strip())
            except:
                return None

        split_coords['Latitude'] =
split_coords['Latitude'].apply(to_float)
        split_coords['Longitude'] =
split_coords['Longitude'].apply(to_float)

```

```

# Keep only valid numeric rows
split_coords = split_coords.dropna(subset=['Latitude',
'Longitude'])

# Add back to the original DataFrame
df_clean = df_valid.loc[split_coords.index].copy()
df_clean['Latitude'] = split_coords['Latitude'].round(5)
df_clean['Longitude'] = split_coords['Longitude'].round(5)

print("Cleaned GPS coordinates:")
print(df_clean[['Vehicle Location', 'Latitude',
'Longitude']].head())

import pandas as pd

# Load CSV
df = pd.read_csv(r"D:\DA\python\Electric_Vehicle_Population_Data
(3).csv")

# Clean Vehicle Location column
df['Vehicle Location'] = df['Vehicle
Location'].astype(str).str.strip()
df = df[df['Vehicle Location'].notna()]
df = df[df['Vehicle Location'] != '']
df = df[df['Vehicle Location'].str.lower() != 'nan']

# Keep only rows with a comma
df_valid = df[df['Vehicle Location'].str.contains(',',
na=False)].copy()

if df_valid.empty:
    print("No valid Vehicle Location entries. df_clean cannot be
created.")
else:
    split_coords = df_valid['Vehicle Location'].str.split(',', n=1,
expand=True)

    if split_coords.shape[1] >= 2:
        split_coords.columns = ['Latitude', 'Longitude']

        def to_float(x):
            try:
                return float(str(x).strip())
            except:
                return None

        split_coords['Latitude'] =
split_coords['Latitude'].apply(to_float)
        split_coords['Longitude'] =

```

```

split_coords['Longitude'].apply(to_float)
    split_coords = split_coords.dropna(subset=['Latitude',
'Longitude'])

    df_clean = df_valid.loc[split_coords.index].copy()
    df_clean['Latitude'] = split_coords['Latitude'].round(5)
    df_clean['Longitude'] = split_coords['Longitude'].round(5)

    # Now df_clean exists, and you can save it
    df_clean.to_csv(r"D:\DA\python\Electric_Vehicle_Population_Clean.csv", index=False)
    print("Cleaned CSV saved successfully.")
else:
    print("Split did not produce two columns. df_clean cannot be created.")

No valid Vehicle Location entries. df_clean cannot be created.

```

Section 2 : Data Exploration Questions

A > What are the top 5 most common EV makes and models in the dataset?

```

import pandas as pd
df = pd.read_csv(r"D:\DA\python\Electric_Vehicle_Population_Data
(3).csv")

df.head()

   VIN (1-10)      County          City State Postal Code Model Year
Make \
0  JTDKN3DP2D     Yakima       Yakima   WA  98902.0  2013
TOYOTA
1  1FMCU0E1XS     Kitsap  Port Orchard   WA  98366.0  2025
FORD
2  JM3KKBHA9R     Kitsap     Kingston   WA  98346.0  2024
MAZDA
3  7SAYGDEE8P Thurston     Olympia   WA  98501.0  2023
TESLA
4  5YJ3E1EB5K Thurston     Rainier   WA  98576.0  2019
TESLA

      Model          Electric Vehicle Type \
0    PRIUS  Plug-in Hybrid Electric Vehicle (PHEV)
1  ESCAPE  Plug-in Hybrid Electric Vehicle (PHEV)
2   CX-90  Plug-in Hybrid Electric Vehicle (PHEV)
3  MODEL Y      Battery Electric Vehicle (BEV)
4  MODEL 3      Battery Electric Vehicle (BEV)

   Clean Alternative Fuel Vehicle (CAFV) Eligibility  Electric
Range \

```

```
0           Not eligible due to low battery range      6.0
1           Clean Alternative Fuel Vehicle Eligible   37.0
2           Not eligible due to low battery range     26.0
3  Eligibility unknown as battery range has not b...    0.0
4           Clean Alternative Fuel Vehicle Eligible 220.0
```

```
Base MSRP  Legislative District  DOL Vehicle ID \
0          0.0                  15.0    165252538
1          0.0                  26.0    278572521
2          0.0                  23.0    275123642
3          0.0                  35.0    249569323
4          0.0                  20.0    283135107
```

Tract	Vehicle Location	Electric Utility	2020 Census
0	POINT (-120.51904 46.59783)	PACIFICORP	5.307700e+10
1	POINT (-122.63847 47.54103)	PUGET SOUND ENERGY INC	5.303509e+10
2	POINT (-122.4977 47.79802)	PUGET SOUND ENERGY INC	5.303509e+10
3	POINT (-122.89165 47.03954)	PUGET SOUND ENERGY INC	5.306701e+10
4	POINT (-122.68993 46.88897)	PUGET SOUND ENERGY INC	5.306701e+10

```
# Top 5 EV Makes
top_makes = df['Make'].value_counts().head(5)
print("Top 5 EV Makes:")
display(top_makes)
```

```
Top 5 EV Makes:
```

```
Make
TESLA        108777
CHEVROLET    18908
NISSAN       16224
FORD          13988
KIA           12849
Name: count, dtype: int64
```

```
# Top 5 EV Models
top_models = df['Model'].value_counts().head(5)
print("Top 5 EV Models:")
display(top_models)
```

```
Top 5 EV Models:
```

```
Model
MODEL Y      54720
MODEL 3       37774
LEAF          13852
MODEL S       7945
BOLT EV       7873
Name: count, dtype: int64
```

B > What is the distribution of EVs by county? Which county has the most registrations?

```
df.columns
Index(['VIN (1-10)', 'County', 'City', 'State', 'Postal Code', 'Model
Year',
       'Make', 'Model', 'Electric Vehicle Type',
       'Clean Alternative Fuel Vehicle (CAFV) Eligibility', 'Electric
Range',
       'Base MSRP', 'Legislative District', 'DOL Vehicle ID',
       'Vehicle Location', 'Electric Utility', '2020 Census Tract'],
      dtype='object')
```

```
# EVs are registered in each county
county_distribution = df['County'].value_counts()

print("Distribution of EVs by County:")
display(county_distribution)
```

Distribution of EVs by County:

```
County
King        130129
Snohomish   32335
Pierce      21624
Clark        15925
Thurston    9506
...
Platte      1
Manatee     1
Escambia    1
Utah        1
Denton      1
Name: count, Length: 236, dtype: int64
```

```
# top county
top_county = county_distribution.idxmax()
top_count = county_distribution.max()
print(f"\nCounty with the most EV registrations: {top_county}
({top_count} vehicles)")
```

County with the most EV registrations: King (130129 vehicles)

C> How has EV adoption changed over different model years?

```
# EVs are registered for each model year
ev_by_year = df['Model Year'].value_counts().sort_index()

print("EV Adoption by Model Year:")
display(ev_by_year)
```

EV Adoption by Model Year:

```
Model Year
2000      8
2002      1
2003      1
2008     20
2010     22
2011    631
2012   1440
2013   4081
2014   3327
2015   4574
2016   5253
2017   8767
2018  14524
2019  11043
2020  12395
2021  20937
2022  29647
2023  60215
2024  49869
2025  29495
2026  5448
Name: count, dtype: int64
```

D > What is the average electric range of EVs in the dataset ?

```
df.columns
Index(['VIN (1-10)', 'County', 'City', 'State', 'Postal Code', 'Model
Year',
       'Make', 'Model', 'Electric Vehicle Type',
       'Clean Alternative Fuel Vehicle (CAFV) Eligibility', 'Electric
Range',
       'Base MSRP', 'Legislative District', 'DOL Vehicle ID',
       'Vehicle Location', 'Electric Utility', '2020 Census Tract'],
      dtype='object')
```

```
# Average Electric Range
average_range = df['Electric Range'].mean()
print("Average Electric Range of EVs:", round(average_range, 2),
      "miles")
```

Average Electric Range of EVs: 42.62 miles

E > What percentage of EVs are eligible for Clean Alternative Fuel Vehicle (CAFV) incentives?

```
total_vehicles = len(df)
eligible_vehicles = df[df['Clean Alternative Fuel Vehicle (CAFV) Eligibility'] == 'Clean Alternative Fuel Vehicle Eligible'].shape[0]

# Calculate percentage
percentage_eligible = (eligible_vehicles / total_vehicles) * 100
print(f"Percentage of EVs eligible for CAFV incentives: {percentage_eligible:.2f}%")
```

Percentage of EVs eligible for CAFV incentives: 29.35%

F > How does the electric range vary across different makes and models?

```
# Average Electric Range by Make
avg_range_by_make = df.groupby('Make')['Electric Range'].mean().sort_values(ascending=False)
print("Average Electric Range by Make:")
display(avg_range_by_make)
```

Average Electric Range by Make:

Make	Avg Electric Range
JAGUAR	181.267606
WHEEGO ELECTRIC CARS	100.000000
TH!NK	100.000000
CHEVROLET	82.355458
FIAT	75.490588
NISSAN	65.670118
SMART	61.686441
AZURE DYNAMICS	56.000000
TESLA	55.961150
PORSCHE	49.788640
LAND ROVER	48.523316
AUDI	38.013809
ALFA ROMEO	33.000000
POLESTAR	32.812458
MITSUBISHI	32.203347
CHRYSLER	32.167143
BENTLEY	31.250000
KIA	30.250370
DODGE	29.045514

BMW	28.391618
TOYOTA	27.806731
MAZDA	25.559971
LINCOLN	24.631579
LAMBORGHINI	22.909091
LEXUS	22.112295
JEEP	21.890893
VOLVO	17.828012
VOLKSWAGEN	16.729367
HONDA	15.838747
MERCEDES - BENZ	15.323832
MINI	14.038972
HYUNDAI	13.090533
FORD	7.579139
FISKER	2.275862
CADILLAC	1.909250
SUBARU	0.628639
BRIGHTDROP	0.000000
ACURA	0.000000
LUCID	0.000000
GMC	0.000000
GENESIS	0.000000
MULLEN AUTOMOTIVE INC.	0.000000
ROLLS-ROYCE	0.000000
RIVIAN	0.000000
RAM	0.000000
VINFEST	0.000000

Name: Electric Range, dtype: float64

```
# Average Electric Range by Make & Model
avg_range_by_model = df.groupby(['Make', 'Model'])['Electric
Range'].mean().sort_values(ascending=False)
print("Average Electric Range by Make and Model:")
# show top 10 models
display(avg_range_by_model.head(10))
```

Average Electric Range by Make and Model:

Make	Model	
TESLA	ROADSTER	230.000000
JAGUAR	I-PACE	181.267606
CHEVROLET	BOLT EV	168.491680
TESLA	MODEL S	165.849339
AUDI	E-TRON	128.594881
TESLA	MODEL X	118.113199
VOLKSWAGEN	E-GOLF	107.096408
PORSCHE	MACAN	105.070632
TOYOTA	RAV4	102.728814
TH!NK	CITY	100.000000

Name: Electric Range, dtype: float64

G > What is the average Base MSRP for each EV model?

```
# Average Base MSRP for Each Model
# Group by Make and Model, then calculate average MSRP
avg_msrp_by_model = df.groupby(['Make', 'Model'])['Base
MSRP'].mean().sort_values(ascending=False)

print("Average Base MSRP for each EV model:")
display(avg_msrp_by_model)
```

Average Base MSRP for each EV model:

Make	Model	Base MSRP
PORSCHE	918	845000.000000
TESLA	ROADSTER	103563.541667
FISKER	KARMA	102000.000000
BMW	740E	90287.037037
CADILLAC	CT6	75095.000000
		...
VOLVO	EX90	0.000000
	V60	0.000000
	S90	0.000000
	S60	0.000000
	XC40	0.000000

Name: Base MSRP, Length: 181, dtype: float64

```
avg_msrp_by_model.head(10)
```

Make	Model	Base MSRP
PORSCHE	918	845000.000000
TESLA	ROADSTER	103563.541667
FISKER	KARMA	102000.000000
BMW	740E	90287.037037
CADILLAC	CT6	75095.000000
BMW	530E	35430.091533
WHEEGO ELECTRIC CARS	WHEEGO	32995.000000
KIA	SOUL	30868.695652
SUBARU	CROSSTREK	24570.957447
MINI	COUNTRYMAN	15601.259446

Name: Base MSRP, dtype: float64

H > Are there any regional trends in EV adoption (e.g., urban vs. rural areas)?

```
# By following steps , can analyze this directly from your EV dataset.
# EVs by County or City
ev_by_region = df['County'].value_counts()
print("EV Registrations by County:")
display(ev_by_region.head(10))
```

EV Registrations by County:

```

County
King           130129
Snohomish      32335
Pierce         21624
Clark          15925
Thurston       9506
Kitsap         8787
Spokane        7370
Whatcom        6406
Benton         3572
Skagit         3067
Name: count, dtype: int64

# Categorize as Urban or Rural
urban_counties = ['King', 'Snohomish', 'Pierce', 'Clark', 'Thurston']
df['Region Type'] = df['County'].apply(lambda x: 'Urban' if x in
urban_counties else 'Rural')

region_summary = df['Region Type'].value_counts(normalize=True) * 100
print("EV Adoption by Region Type:")
display(region_summary)

EV Adoption by Region Type:

Region Type
Urban     80.061368
Rural    19.938632
Name: proportion, dtype: float64

# This means most EVs are registered in urban areas

```

Section 3 : Data Visualization Questions

A > Create a bar chart showing the top 5 EV makes and models by count.

```

# Top 5 EV makes by count
top5_makes = df['Make'].value_counts().head(5)
# Display as a table
print("Top 5 EV Makes by Count:")
display(top5_makes)

```

Top 5 EV Makes by Count:

Make	Count
TESLA	108777
CHEVROLET	18908
NISSAN	16224
FORD	13988
KIA	12849

Name: count, dtype: int64

```
# Top 5 EV Models.
top5_models = df['Model'].value_counts().head(5)

print("Top 5 EV Models by Count:")
display(top5_models)

Top 5 EV Models by Count:

Model
MODEL Y      54720
MODEL 3      37774
LEAF         13852
MODEL S      7945
BOLT EV      7873
Name: count, dtype: int64
```

B > Use a heatmap or choropleth map to visualize EV distribution by county.

```
!pip install plotly

Requirement already satisfied: plotly in c:\users\hp\appdata\local\programs\python\python312\lib\site-packages (6.3.1)
Requirement already satisfied: narwhals>=1.15.1 in c:\users\hp\appdata\local\programs\python\python312\lib\site-packages (from plotly) (2.7.0)
Requirement already satisfied: packaging in c:\users\hp\appdata\local\programs\python\python312\lib\site-packages (from plotly) (25.0)

# Aggregate EV counts by county
ev_by_county = df['County'].value_counts().reset_index()
ev_by_county.columns = ['County', 'EV_Count']
ev_by_county.head()

   County  EV_Count
0    King     130129
1 Snohomish     32335
2   Pierce     21624
3    Clark     15925
4 Thurston      9506

!pip install matplotlib

Collecting matplotlib
  Downloading matplotlib-3.10.6-cp312-cp312-win_amd64.whl.metadata (11 kB)
Collecting contourpy>=1.0.1 (from matplotlib)
  Downloading contourpy-1.3.3-cp312-cp312-win_amd64.whl.metadata (5.5 kB)
Collecting cycler>=0.10 (from matplotlib)
  Downloading cycler-0.12.1-py3-none-any.whl.metadata (3.8 kB)
Collecting fonttools>=4.22.0 (from matplotlib)
```

```
  Downloading fonttools-4.60.1-cp312-cp312-win_amd64.whl.metadata (114 kB)
Collecting kiwisolver>=1.3.1 (from matplotlib)
  Downloading kiwisolver-1.4.9-cp312-cp312-win_amd64.whl.metadata (6.4 kB)
Requirement already satisfied: numpy>=1.23 in c:\users\hp\appdata\local\programs\python\python312\lib\site-packages (from matplotlib) (2.3.3)
Requirement already satisfied: packaging>=20.0 in c:\users\hp\appdata\local\programs\python\python312\lib\site-packages (from matplotlib) (25.0)
Collecting pillow>=8 (from matplotlib)
  Downloading pillow-11.3.0-cp312-cp312-win_amd64.whl.metadata (9.2 kB)
Collecting pyparsing>=2.3.1 (from matplotlib)
  Downloading pyparsing-3.2.5-py3-none-any.whl.metadata (5.0 kB)
Requirement already satisfied: python-dateutil>=2.7 in c:\users\hp\appdata\local\programs\python\python312\lib\site-packages (from matplotlib) (2.9.0.post0)
Requirement already satisfied: six>=1.5 in c:\users\hp\appdata\local\programs\python\python312\lib\site-packages (from python-dateutil>=2.7->matplotlib) (1.17.0)
  Downloading matplotlib-3.10.6-cp312-cp312-win_amd64.whl (8.1 MB)
    0.0/8.1 MB ? eta ----- 0.8/8.1 MB 5.6 MB/s eta
0:00:02
    2.1/8.1 MB 5.9 MB/s eta
0:00:02
    3.7/8.1 MB 6.6 MB/s eta
0:00:01
    4.7/8.1 MB 6.8 MB/s eta
0:00:01
    6.0/8.1 MB 6.3 MB/s eta
0:00:01
    7.6/8.1 MB 6.3 MB/s eta
0:00:01
    8.1/8.1 MB 6.2 MB/s
0:00:01
  Downloading contourpy-1.3.3-cp312-cp312-win_amd64.whl (226 kB)
  Downloading cycler-0.12.1-py3-none-any.whl (8.3 kB)
  Downloading fonttools-4.60.1-cp312-cp312-win_amd64.whl (2.3 MB)
    0.0/2.3 MB ? eta ----- 0.8/2.3 MB 3.0 MB/s eta
0:00:01
    2.1/2.3 MB 4.5 MB/s eta
0:00:01
    2.3/2.3 MB 4.4 MB/s
0:00:00
  Downloading kiwisolver-1.4.9-cp312-cp312-win_amd64.whl (73 kB)
```

```
Downloading pillow-11.3.0-cp312-cp312-win_amd64.whl (7.0 MB)
----- 0.0/7.0 MB ? eta ----- 0.5/7.0 MB 4.2 MB/s eta
0:00:02 ----- 1.6/7.0 MB 4.2 MB/s eta
0:00:02 ----- 2.6/7.0 MB 4.6 MB/s eta
0:00:01 ----- 3.1/7.0 MB 4.6 MB/s eta
0:00:01 ----- 4.2/7.0 MB 4.3 MB/s eta
0:00:01 ----- 5.2/7.0 MB 4.4 MB/s eta
0:00:01 ----- 6.3/7.0 MB 4.5 MB/s eta
0:00:01 ----- 7.0/7.0 MB 4.4 MB/s
0:00:01
Downloading pyparsing-3.2.5-py3-none-any.whl (113 kB)
Installing collected packages: pyparsing, pillow, kiwisolver,
fonttools, cycler, contourpy, matplotlib
----- 0/7 [pyparsing]
----- 0/7 [pyparsing]
----- 0/7 [pyparsing]
----- 1/7 [pillow]
----- 2/7 [kiwisolver]
----- 3/7 [fonttools]
```



```
Successfully installed contourpy-1.3.3 cycler-0.12.1 fonttools-4.60.1  
kiwisolver-1.4.9 matplotlib-3.10.6 pillow-11.3.0 pyparsing-3.2.5
```

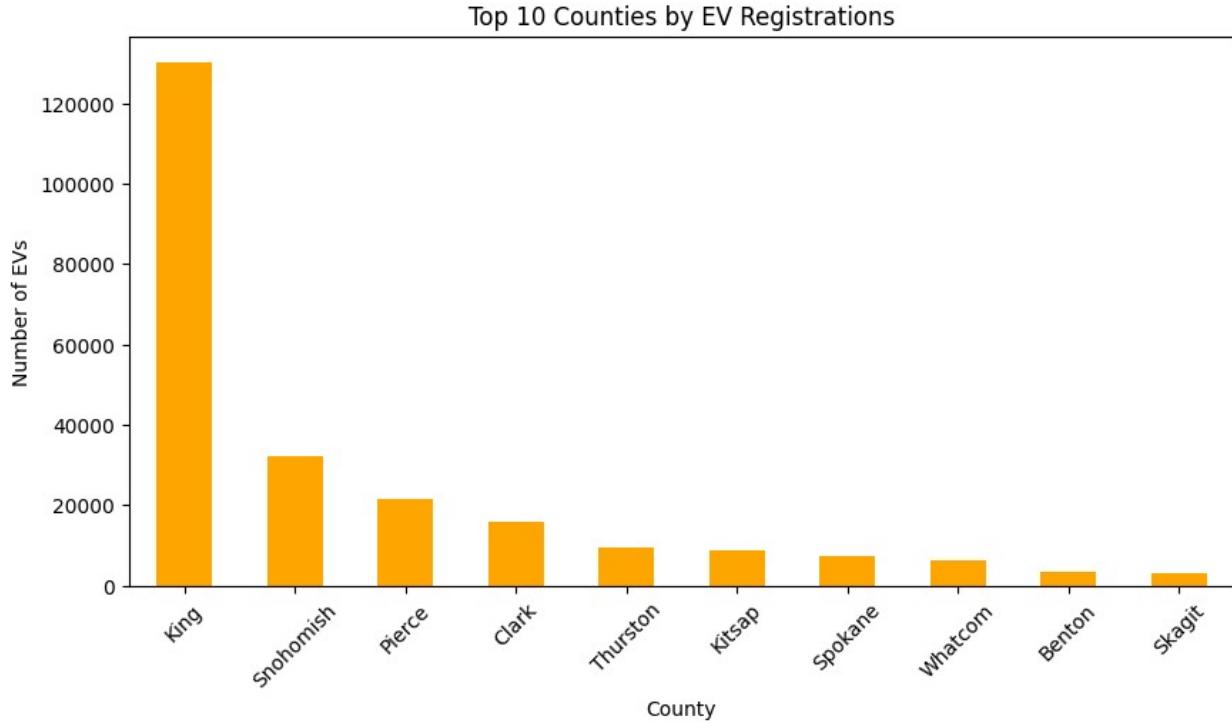
```
import pandas as pd
import matplotlib.pyplot as plt

# Aggregate EV counts by county
ev_by_county = df['County'].value_counts().reset_index()
ev_by_county.columns = ['County', 'EV Count']
```

```
# Plot top 10 counties
top_counties = ev_by_county.head(10)

top_counties.plot(kind='bar', x='County', y='EV_Count',
figsize=(10,5), color='orange', legend=False)
plt.title("Top 10 Counties by EV Registrations")
plt.xlabel("County")
plt.ylabel("Number of EVs")
plt.xticks(rotation=45)
plt.show()

Matplotlib is building the font cache; this may take a moment.
```



C > Create a line graph showing the trend of EV adoption by model year.

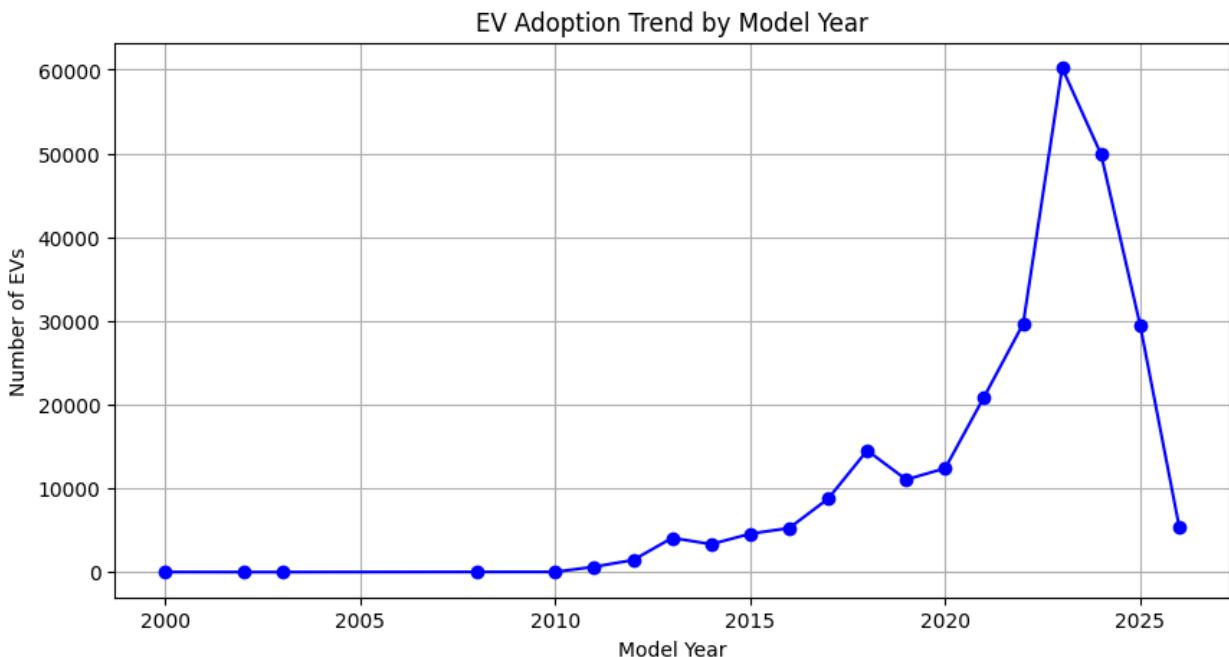
```
# EVs by Model Year
ev_by_year = df['Model Year'].value_counts().sort_index()

print(ev_by_year)

Model Year
2000      8
2002      1
2003      1
2008     20
2010     22
2011   631
```

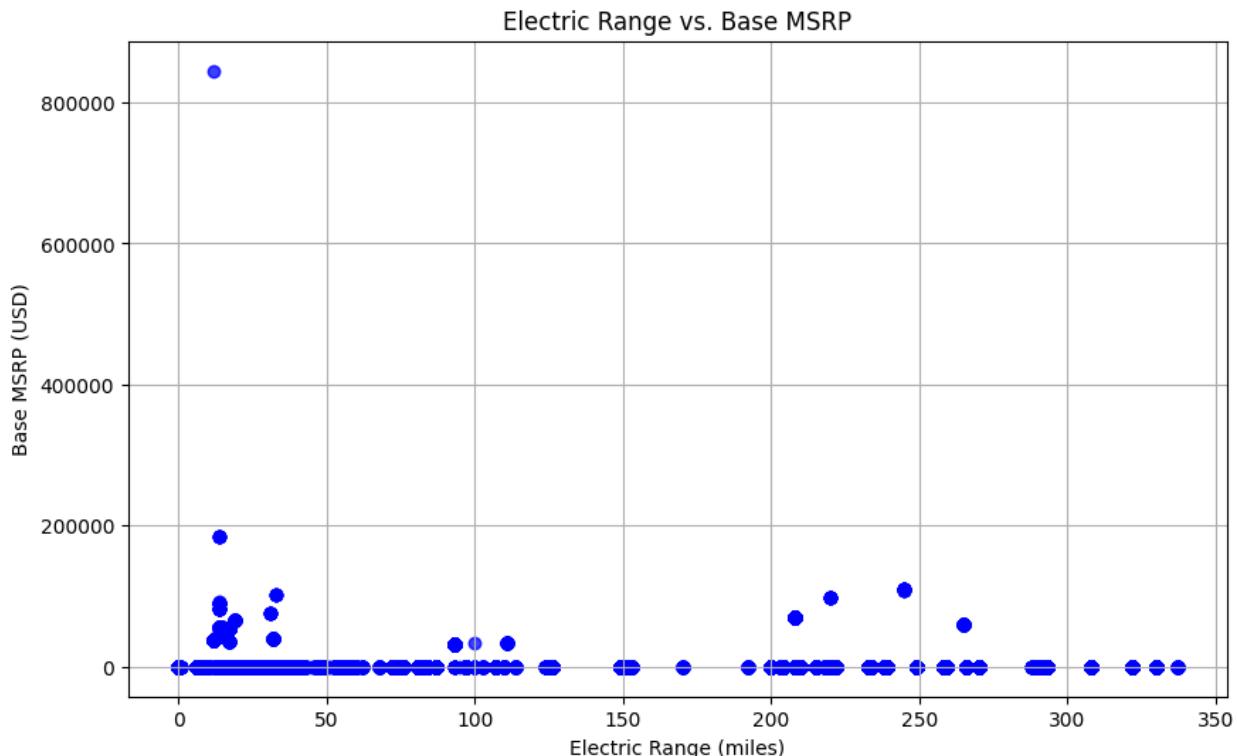
```
2012    1440
2013    4081
2014    3327
2015    4574
2016    5253
2017    8767
2018   14524
2019   11043
2020   12395
2021   20937
2022   29647
2023   60215
2024   49869
2025   29495
2026    5448
Name: count, dtype: int64
```

```
# Creating the line graph
plt.figure(figsize=(10,5))
plt.plot(ev_by_year.index, ev_by_year.values, marker='o',
color='blue')
plt.title("EV Adoption Trend by Model Year")
plt.xlabel("Model Year")
plt.ylabel("Number of EVs")
plt.grid(True)
plt.show()
```



D > Generate a scatter plot comparing electric range vs. base MSRP to see pricing trends

```
# Create a scatter plot
plt.figure(figsize=(10,6))
plt.scatter(df['Electric Range'], df['Base MSRP'], alpha=0.5,
color='blue')
plt.title("Electric Range vs. Base MSRP")
plt.xlabel("Electric Range (miles)")
plt.ylabel("Base MSRP (USD)")
plt.grid(True)
plt.show()
```



E > Plot a pie chart showing the proportion of CAFV-eligible vs. non-eligible EVs.

```
# Count CAFV eligibility
cafv_counts = df['Clean Alternative Fuel Vehicle (CAFV) Eligibility'].value_counts()
cafv_counts

Clean Alternative Fuel Vehicle (CAFV) Eligibility
Eligibility unknown as battery range has not been researched    160888
Clean Alternative Fuel Vehicle Eligible                      76819
Not eligible due to low battery range                     23991
Name: count, dtype: int64

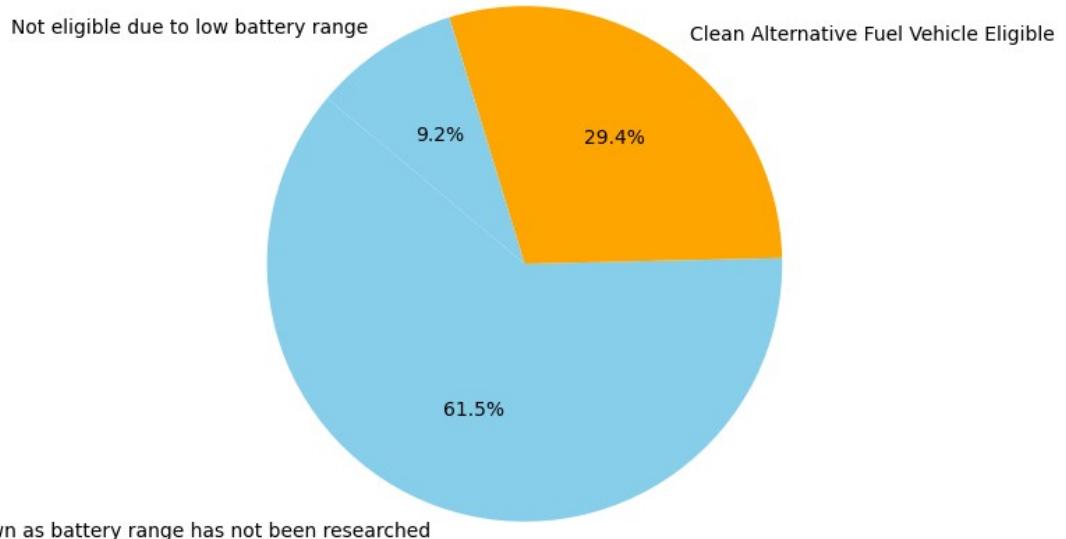
# Plot pie chart
plt.figure(figsize=(6,6))
plt.pie(cafv_counts, labels=cafv_counts.index, autopct='%.1f%%',
```

```

startangle=140, colors=['skyblue', 'orange'])
plt.title("Proportion of CAFV-Eligible vs Non-Eligible EVs")
plt.show()

```

Proportion of CAFV-Eligible vs Non-Eligible EVs



F > Use a geospatial map to display EV registrations based on vehicle location.

```

df.columns
Index(['VIN (1-10)', 'County', 'City', 'State', 'Postal Code', 'Model
Year',
       'Make', 'Model', 'Electric Vehicle Type',
       'Clean Alternative Fuel Vehicle (CAFV) Eligibility', 'Electric
Range',
       'Base MSRP', 'Legislative District', 'DOL Vehicle ID',
       'Vehicle Location', 'Electric Utility', '2020 Census Tract',
       'Region Type'],
      dtype='object')

# Count EVs by county
ev_by_county = df['County'].value_counts().reset_index()
ev_by_county.columns = ['County', 'EV_Count']

# Display top 10 counties
ev_by_county.head(10)

   County  EV_Count
0    King     130129
1 Snohomish      32335

```

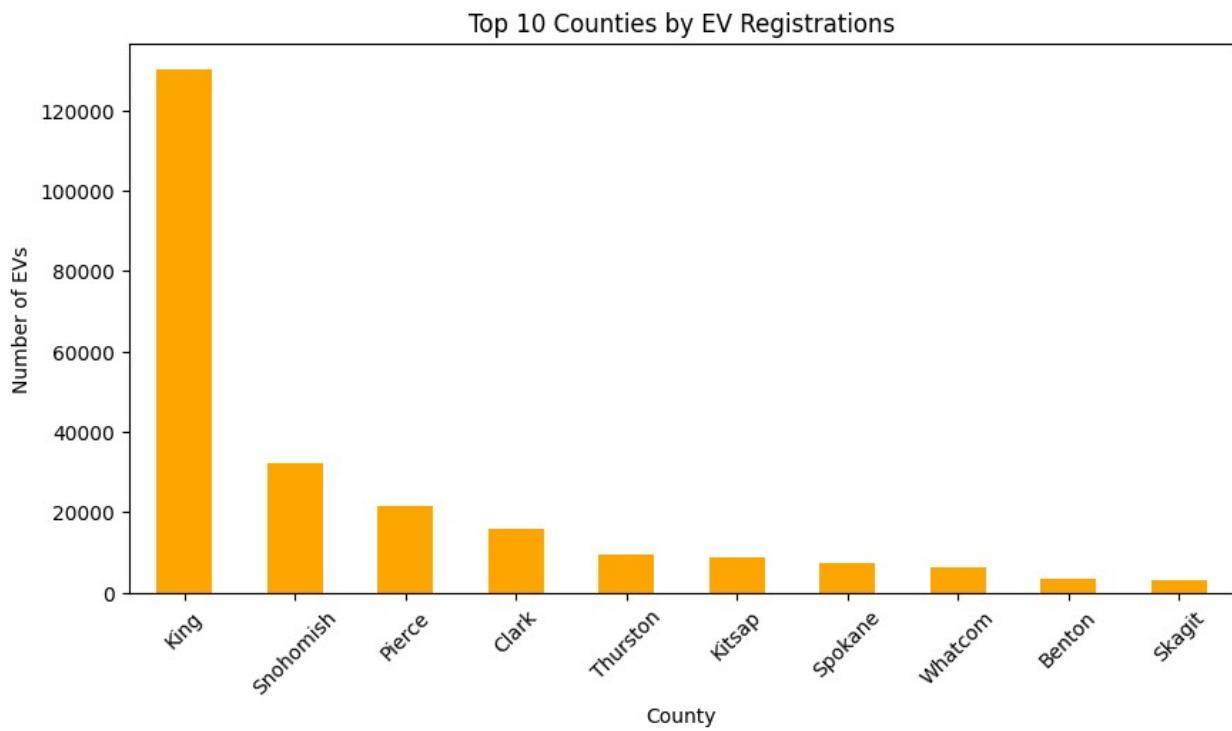
```

2      Pierce    21624
3      Clark     15925
4    Thurston    9506
5      Kitsap    8787
6      Spokane    7370
7    Whatcom    6406
8      Benton    3572
9      Skagit    3067

# Visualize with a bar chart
import matplotlib.pyplot as plt

top_counties = ev_by_county.head(10)
top_counties.plot(kind='bar', x='County', y='EV_Count',
figsize=(10,5), color='orange', legend=False)
plt.title("Top 10 Counties by EV Registrations")
plt.xlabel("County")
plt.ylabel("Number of EVs")
plt.xticks(rotation=45)
plt.show()

```



```
# data set not having latitude and longitude columns
```

Section 4 : Linear Regression Model Questions

A > How can we use Linear Regression to predict the Electric Range of a vehicle?

Ans : Linear Regression predicts a continuous target variable (here, Electric Range) based on one or more features (independent variables) like: Base MSRP (price of the vehicle), Battery Capacity, Vehicle Weight, Motor Power, Model Year

```
!pip install scikit-learn

Requirement already satisfied: scikit-learn in c:\users\hp\appdata\local\programs\python\python312\lib\site-packages (1.7.2)
Requirement already satisfied: numpy>=1.22.0 in c:\users\hp\appdata\local\programs\python\python312\lib\site-packages (from scikit-learn) (2.3.3)
Requirement already satisfied: scipy>=1.8.0 in c:\users\hp\appdata\local\programs\python\python312\lib\site-packages (from scikit-learn) (1.16.2)
Requirement already satisfied: joblib>=1.2.0 in c:\users\hp\appdata\local\programs\python\python312\lib\site-packages (from scikit-learn) (1.5.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in c:\users\hp\appdata\local\programs\python\python312\lib\site-packages (from scikit-learn) (3.6.0)

import sklearn
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

print("scikit-learn version:", sklearn.__version__)

scikit-learn version: 1.7.2

df.columns

Index(['VIN (1-10)', 'County', 'City', 'State', 'Postal Code', 'Model Year',
       'Make', 'Model', 'Electric Vehicle Type',
       'Clean Alternative Fuel Vehicle (CAFV) Eligibility', 'Electric Range',
       'Base MSRP', 'Legislative District', 'DOL Vehicle ID',
       'Vehicle Location', 'Electric Utility', '2020 Census Tract'],
      dtype='object')

# numeric columns in your dataset
print(df.select_dtypes(include=['int64', 'float64']).columns)

Index(['Postal Code', 'Model Year', 'Electric Range', 'Base MSRP',
       'Legislative District', 'DOL Vehicle ID', '2020 Census Tract'],
      dtype='object')

# Convert 'Make' and 'Model' to numeric features
df_encoded = pd.get_dummies(df, columns=['Make', 'Model'],
                           drop_first=True)
```

```
# Now df_encoded has numeric columns instead of text for Make and Model
```

```
print(df_encoded.head())
```

Year \	VIN (1-10)	County	City	State	Postal Code	Model
0	JTDKN3DP2D	Yakima	Yakima	WA	98902.0	2013
1	1FMCU0E1XS	Kitsap	Port Orchard	WA	98366.0	2025
2	JM3KKBHA9R	Kitsap	Kingston	WA	98346.0	2024
3	7SAYGDEE8P	Thurston	Olympia	WA	98501.0	2023
4	5YJ3E1EB5K	Thurston	Rainier	WA	98576.0	2019

	Electric Vehicle Type \
0	Plug-in Hybrid Electric Vehicle (PHEV)
1	Plug-in Hybrid Electric Vehicle (PHEV)
2	Plug-in Hybrid Electric Vehicle (PHEV)
3	Battery Electric Vehicle (BEV)
4	Battery Electric Vehicle (BEV)

Range \	Clean Alternative Fuel Vehicle (CAFV) Eligibility	Electric
0	Not eligible due to low battery range	6.0
1	Clean Alternative Fuel Vehicle Eligible	37.0
2	Not eligible due to low battery range	26.0
3	Eligibility unknown as battery range has not b...	0.0
4	Clean Alternative Fuel Vehicle Eligible	220.0

Model_XC40 \	Base MSRP	...	Model_WHEEGO	Model_WRANGLER	Model_X3	Model_X5
0	0.0	...	False	False	False	False
1	0.0	...	False	False	False	False
2	0.0	...	False	False	False	False
3	0.0	...	False	False	False	False
4	0.0	...	False	False	False	False

Model_XC60	Model_XC90	Model_XM	Model_ZDX	Model_ZEV0
------------	------------	----------	-----------	------------

```

0      False     False     False     False     False     False
1      False     False     False     False     False     False
2      False     False     False     False     False     False
3      False     False     False     False     False     False
4      False     False     False     False     False     False

[5 rows x 240 columns]

# Numeric features
numeric_features = ['Base MSRP', 'Model Year'] # replace with numeric
# columns you have

# Combine with encoded categorical features
X = pd.concat([df[numeric_features], df_encoded], axis=1)
y = df['Electric Range']

# Train-test split
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Automatically select numeric columns
numeric_features = df.select_dtypes(include=['int64',
'float64']).columns.tolist()

# Make sure target 'Electric Range' is excluded from features
if 'Electric Range' in numeric_features:
    numeric_features.remove('Electric Range')

X = df[numeric_features] # only numeric features
y = df['Electric Range']

# Handle missing numeric values
X = X.fillna(X.mean()) # fill missing numeric values with mean
y = y.fillna(y.mean())

# Train-test split and Linear Regression
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

model = LinearRegression()
model.fit(X_train, y_train)

LinearRegression()

# Evaluate Model
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

```

```

print("Mean Squared Error:", mse)
print("R-squared:", r2)

Mean Squared Error: 4636.422413685628
R-squared: 0.3006988849618796

```

B > What independent variables (features) can be used to predict Electric Range? (e.g., Model Year, Base MSRP, Make)

Ans : For predicting the Electric Range of a vehicle, you should use features that influence how far a car can travel on a single charge. These features are called independent variables in regression analysis. Like Price & Market Variables : Base MSRP → Often correlates with battery size and advanced technology. Trim Level → Higher trims may have better motors or larger batteries. Make & Model : Make (Manufacturer) → Different manufacturers design vehicles with different efficiencies. Model → Specific model design affects energy usage. Technology & Year : Model Year → Newer models usually have better battery tech and efficiency. Technology Features → Regenerative braking, energy-saving modes, etc For Ex :

```

df.head()

      VIN (1-10)    County          City State Postal Code Model Year
Make \
0   JTDKN3DP2D    Yakima        Yakima  WA     98902.0  2013
41
1   1FMCU0E1XS    Kitsap  Port Orchard  WA     98366.0  2025
13
2   JM3KKBHA9R    Kitsap        Kingston  WA     98346.0  2024
26
3   7SAYGDEE8P  Thurston       Olympia  WA     98501.0  2023
39
4   5YJ3E1EB5K  Thurston      Rainier  WA     98576.0  2019
39

      Model          Electric Vehicle Type \
0     118  Plug-in Hybrid Electric Vehicle (PHEV)
1     55  Plug-in Hybrid Electric Vehicle (PHEV)
2     39  Plug-in Hybrid Electric Vehicle (PHEV)
3    105      Battery Electric Vehicle (BEV)
4    102      Battery Electric Vehicle (BEV)

      Clean Alternative Fuel Vehicle (CAFV) Eligibility  Electric
Range \
0           Not eligible due to low battery range      6.0
1           Clean Alternative Fuel Vehicle Eligible    37.0
2           Not eligible due to low battery range     26.0
3  Eligibility unknown as battery range has not b...     0.0

```

```

4           Clean Alternative Fuel Vehicle Eligible      220.0

   Base MSRP  Legislative District  DOL Vehicle ID \
0       0.0          15.0    165252538
1       0.0          26.0    278572521
2       0.0          23.0    275123642
3       0.0          35.0    249569323
4       0.0          20.0    283135107

          Vehicle Location      Electric Utility 2020 Census
Tract
0 POINT (-120.51904 46.59783)          PACIFICORP
5.307700e+10
1 POINT (-122.63847 47.54103)  PUGET SOUND ENERGY INC
5.303509e+10
2 POINT (-122.4977 47.79802)  PUGET SOUND ENERGY INC
5.303509e+10
3 POINT (-122.89165 47.03954)  PUGET SOUND ENERGY INC
5.306701e+10
4 POINT (-122.68993 46.88897)  PUGET SOUND ENERGY INC
5.306701e+10

# Features available in your dataset
features = ['Base MSRP', 'Model Year', 'Make', 'Model']

# Target variable
target = 'Electric Range'

# Create feature and target DataFrames
X = df[features]
y = df[target]

X = X.copy()
X['Base MSRP'] = X['Base MSRP'].fillna(X['Base MSRP'].mean())
X['Model Year'] = X['Model Year'].fillna(X['Model Year'].mean())
X[['Make', 'Model']] = X[['Make', 'Model']].fillna('Unknown')

# Encode Categorical Variables
X_encoded = pd.get_dummies(X, columns=['Make', 'Model'],
drop_first=True)
# Split Data
X_train, X_test, y_train, y_test = train_test_split(X_encoded, y,
test_size=0.2, random_state=42)

# Train Linear Regression
model = LinearRegression()
model.fit(X_train, y_train)

LinearRegression()

```

```

# Evaluate Model
y_pred = model.predict(X_test)

mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print("Mean Squared Error:", mse)
print("R2 Score:", r2)

Mean Squared Error: 3181.4734499626393
R2 Score: 0.5201455491941502

# Interpret Base MSRP Influence
coeff_df = pd.DataFrame({'Feature': X_encoded.columns, 'Coefficient': model.coef_})
print(coeff_df.sort_values(by='Coefficient', ascending=False).head(10))

      Feature  Coefficient
55    Model_9    436.590933
101   Model_55   106.770988
69    Model_23   101.935617
34    Make_33    101.602689
50    Model_4     91.082590
145   Model_99    77.466139
108   Model_62    76.620076
180   Model_134   74.910094
89    Model_43    70.739262
207   Model_161   67.073938

```

C > How do we handle categorical variables like Make and Model in regression analysis?

Ans : In regression analysis, categorical variables like Make and Model cannot be used directly, because regression algorithms require numerical inputs. We need to convert these categorical variables into numbers

```

import pandas as pd

# One-hot encode 'Make' and 'Model'
X_encoded = pd.get_dummies(X, columns=['Make', 'Model'],
drop_first=True)

```

Use One-Hot Encoding for Make and Model.If there are hundreds of categories, consider keeping top N frequent categories and label others as "Other". For ex:

```

# Original features
features = ['Base MSRP', 'Model Year', 'Make', 'Model']
X = df[features]

# One-hot encode categorical variables

```

```

X_encoded = pd.get_dummies(X, columns=['Make', 'Model'],
drop_first=True)

# Now X_encoded can be used for Linear Regression

X = df[['Base MSRP', 'Model Year', 'Make', 'Model']]
y = df['Electric Range']

#Handle Missing Values
X = X.copy()
X['Base MSRP'] = X['Base MSRP'].fillna(X['Base MSRP'].mean())
X['Model Year'] = X['Model Year'].fillna(X['Model Year'].mean())
X[['Make', 'Model']] = X[['Make', 'Model']].fillna('Unknown')
y = y.fillna(y.mean())

# Convert Categorical Columns
X_encoded = pd.get_dummies(X, columns=['Make', 'Model'],
drop_first=True)
#Split Data
X_train, X_test, y_train, y_test = train_test_split(X_encoded, y,
test_size=0.2, random_state=42)

# Train Model
model = LinearRegression()
model.fit(X_train, y_train)
# Evaluate Model
y_pred = model.predict(X_test)

print("R2 Score:", r2_score(y_test, y_pred))
print("Mean Squared Error:", mean_squared_error(y_test, y_pred))

R2 Score: 0.5201455491941502
Mean Squared Error: 3181.4734499626393

#Predict for a New EV
new_vehicle = pd.DataFrame({
    'Base MSRP':[55000],
    'Model Year':[2022],
    'Make':['TESLA'],
    'Model':['Model 3']
})

new_vehicle_encoded = pd.get_dummies(new_vehicle)
new_vehicle_encoded =
new_vehicle_encoded.reindex(columns=X_encoded.columns, fill_value=0)

predicted_range = model.predict(new_vehicle_encoded)
print("Predicted Electric Range:", predicted_range[0])

Predicted Electric Range: -19.010836518660653

```

D > What is the R² score of the model, and what does it indicate about prediction accuracy?

Ans : R² (R-squared) is a statistical measure of how well your regression model explains the variation in the target variable. It ranges from 0 to 1 (sometimes slightly negative if the model is very bad). R² = 1 → Perfect prediction (all points fit the model exactly) R² = 0 → Model does not explain any variation in the data R² tells us how much of the variation in Electric Range can be explained by our model using Base MSRP, Model Year, Make, and Model. For ex.

```
# Get R2 Score in Python
from sklearn.metrics import r2_score

r2 = r2_score(y_test, y_pred)
print("R2 Score:", r2)

R2 Score: 0.5201455491941502
```

E > How does the Base MSRP influence the Electric Range according to the regression model?

In Linear Regression, each feature (like Base MSRP) gets a coefficient. The coefficient shows how much the target (Electric Range) changes when that feature changes — keeping others constant. So if the Base MSRP coefficient is positive, it means: As the Base MSRP (price) increases, the Electric Range tends to increase. If it's negative, then: As the Base MSRP increases, the Electric Range tends to decrease.

```
# Check it in Python
coeff_df = pd.DataFrame({'Feature': X_encoded.columns, 'Coefficient': model.coef_})
coeff_df[coeff_df['Feature'] == 'Base MSRP']

      Feature  Coefficient
0  Base MSRP     -0.000804
```

For ex : If the coefficient of Base MSRP is -0.000804 , it means: For every 1 unit increase in Base MSRP, the Electric Range decreases slightly by about 0.000804 miles. Since the coefficient is negative, it shows a small negative relationship between vehicle price and range.

The coefficient of Base MSRP is -0.000804, which means there is a slight negative relationship between vehicle price and electric range. As the Base MSRP increases, the electric range decreases a little. This may be because higher-priced EVs often focus on premium features or performance rather than only extending range.

E > What steps are needed to improve the accuracy of the Linear Regression model?

1 > Add More Relevant Features Right now, the model only uses Base MSRP, Model Year, Make, and Model

2 > Remove or Handle Outliers Outliers (unusual values) can mislead the model. You can detect them using:

```
df.describe()
```

	Postal Code	Model Year	Make	Model	\
count	261688.000000	261698.000000	261698.000000	261698.000000	
mean	98176.150699	2021.772493	28.929549	99.273812	
std	2555.753410	3.034041	13.125556	36.812807	
min	1469.000000	2000.000000	0.000000	0.000000	
25%	98052.000000	2020.000000	17.000000	90.000000	
50%	98133.000000	2023.000000	39.000000	103.000000	
75%	98382.000000	2024.000000	39.000000	106.000000	
max	99577.000000	2026.000000	45.000000	180.000000	
Vehicle ID	Electric Range	Base MSRP	Legislative District	DOL	
count	261695.000000	261698.000000	261070.000000		
mean	42.615071	695.503563	28.881955		
std	81.226054	6942.979857	14.889697		
min	0.000000	0.000000	1.000000		
25%	0.000000	0.000000	17.000000		
50%	0.000000	0.000000	32.000000		
75%	35.000000	0.000000	42.000000		
max	337.000000	845000.000000	49.000000		
	2020 Census Tract				
count	2.616880e+05				
mean	5.297261e+10				
std	1.628791e+09				
min	1.001020e+09				
25%	5.303301e+10				
50%	5.303303e+10				
75%	5.305307e+10				
max	6.601095e+10				

3 > Normalize / Scale Numeric Data Large differences in feature scales can confuse the model.
Use:

```
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X_encoded)
```

This ensures all features have equal importance.

4 > Encode Categorical Data Properly Use One-Hot Encoding (as we did) to convert Make and Model correctly. If there are too many categories, group the least common ones into "Other"

5 > Remove Irrelevant or Correlated Features Too many similar features can reduce performance.

6 > Try Advanced Models Sometimes Linear Regression isn't enough

7 > Cross-Validation Instead of using a single train-test split, use:

```
from sklearn.model_selection import cross_val_score
scores = cross_val_score(model, X_encoded, y, cv=5)
print(scores.mean())
0.5232756907684915
```

F > Can we use this model to predict the range of new EV models based on their specifications?

Yes, We can use the Linear Regression model to predict the Electric Range of new EV models based on their specifications — like Base MSRP, Model Year, Make, and Model. For Ex :

```
# Example: Predict Electric Range for a new EV
new_vehicle = pd.DataFrame({
    'Base MSRP': [55000],
    'Model Year': [2022],
    'Make': ['TESLA'],
    'Model': ['Model 3']
})

# Convert to same format as training data
new_vehicle_encoded = pd.get_dummies(new_vehicle)
new_vehicle_encoded =
new_vehicle_encoded.reindex(columns=X_encoded.columns, fill_value=0)

# Predict
predicted_range = model.predict(new_vehicle_encoded)
print("Predicted Electric Range:", predicted_range[0], "miles")

Predicted Electric Range: -19.010836518660653 miles
```

Conclusion : This analysis of data helps us understand which electric vehicles are most popular, how EV adoption varies by region, and how incentives affect registrations. Cleaning the data made it easier to work with, and visualizations showed clear trends. Using a regression model, we found that factors like Base MSRP, Model Year, and Make can help predict a vehicle's electric range.